Music classification & generation with spectrograms

Background: Sound can have different digital representations: usually it is stored by encoding the shape of the waveform as it changes over time, but for analysis we often make use of visually inspectable spectrograms - obtained from a waveform by computing the Fourier transform of overlapping windows of the audio signal, and stacking the results into a 2D array. These spectrograms exhibit a lot of structure, and modelling it can enable sound - in this case music - classification, generation, even recommendation e.g. speech recognition or music recommendation.

Project setup: We provide a notebook with loading the GTZAN dataset (containing the sound files and genre labels) and the spectrograms, and training a simple CNN on them. One can use other datasets as well (e.g. FMA-small dataset).

Project map:

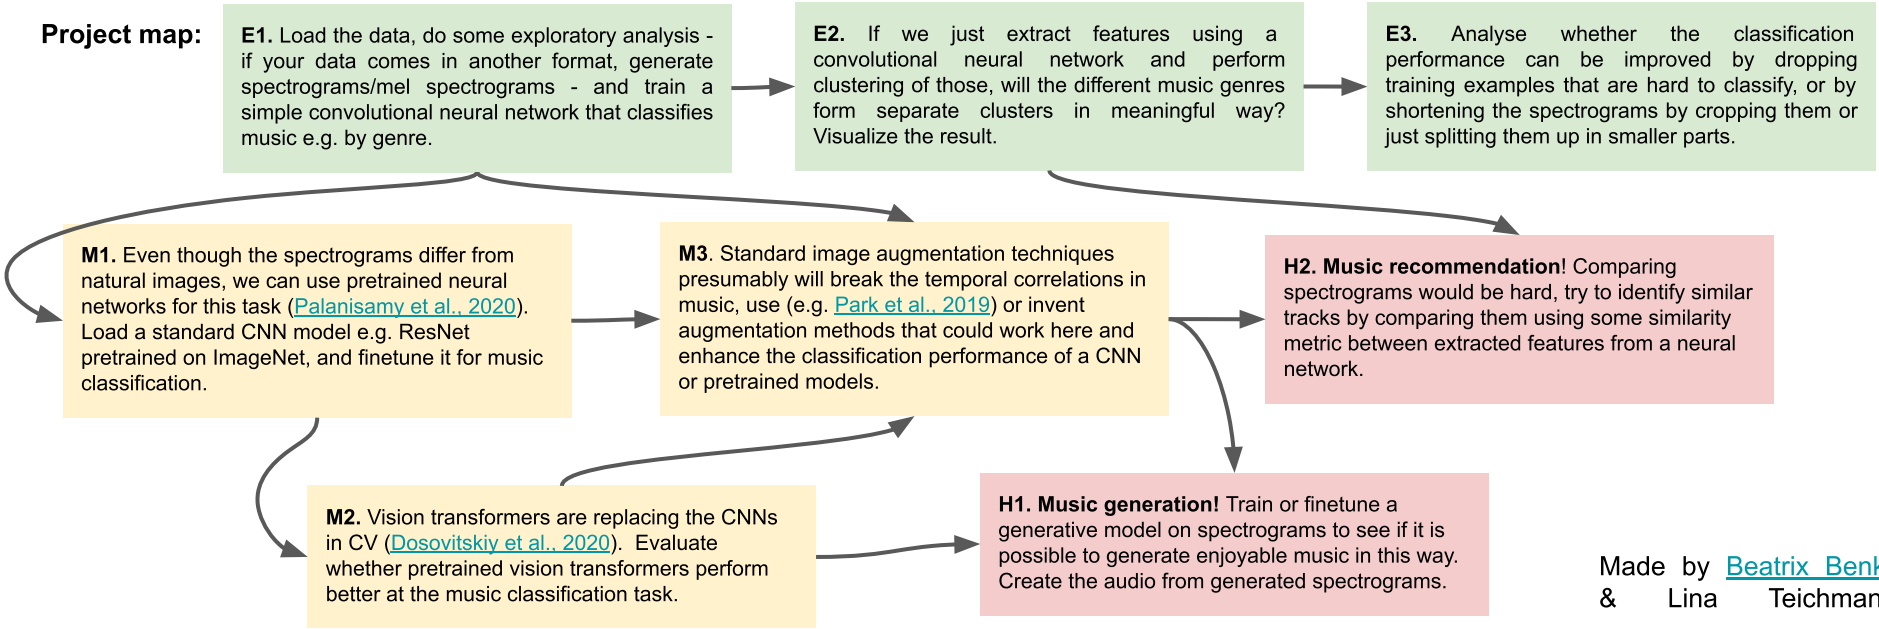

Made by Beatrix Benkő & Lina Teichmann